

Informed search

- Use domain knowledge!

- Are we getting close to the goal?
- Use a heuristic function that estimates how close a state is to the goal
- A heuristic does NOT have to be perfect!
- Example of strategies:
 1. Greedy best-first search
 2. A* search
 3. IDA*

使用领域知识

启发函数不需要很完美

Greedy search

- Evaluation function $h(n)$ (heuristic)
- $h(n)$ estimates the cost from n to the goal
- Example: $h_{SLD}(n)$ = straight-line distance from n to Sault Ste Marie
- Greedy search expands the node that **appears** to be closest to goal

Greedy search : Pseudo-code

```
function GREEDY-BEST-FIRST-SEARCH(initialState, goalTest)
    returns SUCCESS or FAILURE : /* Cost  $f(n) = h(n)$  */

    frontier = Heap.new(initialState)
    explored = Set.new()

    while not frontier.isEmpty():
        state = frontier.deleteMin()
        explored.add(state)

        if goalTest(state):
            return SUCCESS(state)

        for neighbor in state.neighbors():
            if neighbor not in frontier ∪ explored:
                frontier.insert(neighbor)
            else if neighbor in frontier:
                frontier.decreaseKey(neighbor) → 没什么用,
                                                因为  $h(n)$  是固定的

    return FAILURE
```

A* search

- Minimize the total estimated solution cost
- Combines:
 - $g(n)$: cost to reach node n
 - $h(n)$: cost to get from n to the goal
 - $f(n) = g(n) + h(n)$

$f(n)$ is the estimated cost of the cheapest solution through n

有一个中转点

A* search - Pseudo-code

```
function A-STAR-SEARCH(initialState, goalTest)
    returns SUCCESS or FAILURE : /* Cost  $f(n) = g(n) + h(n)$  */
    frontier = Heap.new(initialState)
    explored = Set.new()

    while not frontier.isEmpty():
        state = frontier.deleteMin()
        explored.add(state)

        if goalTest(state):
            return SUCCESS(state)

        for neighbor in state.neighbors():
            if neighbor not in frontier ∪ explored:
                frontier.insert(neighbor)
            else if neighbor in frontier:
                frontier.decreaseKey(neighbor)

    return FAILURE
```

与 Greedy search 相比，仅更改了 cost 函数。

Admissible heuristics

- An **admissible** heuristic never overestimates the cost to reach the goal, that is it is **optimistic**
- A heuristic h is admissible if
$$\forall \text{node } n, h(n) \leq h^*(n)$$
where h^* is true cost to reach the goal from n .
- h_{SLD} (used as a heuristic in the map example) is admissible because it is by definition the shortest distance (straight line) between two points.

A* Optimality

If $h(n)$ is admissible, A* using tree search is optimal.

Rationale:

- Suppose G_o is the optimal goal.
Suppose G_s is some suboptimal goal.
Suppose n is on the shortest path to G_o .
 - $f(G_o) = g(G_o)$ since $h(G_o) = 0$
 - $f(G_s) = g(G_s)$ since $h(G_s) = 0$
 - $f(G_s) > f(G_o)$ since G_s is suboptimal
- Then $f(G_o) > f(G_s) \dots (1)$
- $h(n) \leq h^*(n)$ since h is admissible
 $g(n) + h(n) \leq g(n) + h^*(n) = g(G_o) = f(G_o)$
- Then, $f(n) \leq f(G_o) \dots (2)$

G 指的是到 goal 的序列

From (1) and (2) $f(G_s) > f(n)$, so
A* will never select G_s during the search and hence A* is optimal.



We always will go toward G_o rather than to go G_s .

25

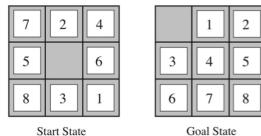
A* - PF Metrics

- Complete:** Yes.
- Time:** exponential
- Space:** keeps every node in memory, the biggest problem
- Optimal:** Yes!

Search Efficiency of Heuristics

Heuristics for 8-puzzle

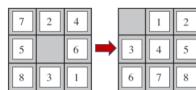
- The solution is 26 steps long.
- $h_1(n)$ = number of misplaced tiles
- $h_1(n) = 8$
- $h_2(n)$ = total Manhattan distance (sum of the horizontal and vertical distances).
- Tiles 1 to 8 in the start state gives: $h_2 = 3+1+2+2+2+3+3+2 = 18$ which does not overestimate the true solution.



- $h_{mis}(s) = \# \text{misplaced tiles} \in [0,8]$: **Admissible**.
- $h_{1stp}(s) = \#(1\text{-step move})$ to reach the goal configuration: **Admissible**.

➤ $h_{1stp}(s) \geq h_{mis}(s) \Rightarrow h_{1stp}(s)$ is '**better**' than $h_{mis}(s)$.

What does '**better**' mean?



Dominance

- For **admissible** h_1 and h_2 , if $h_1(s) \geq h_2(s)$ for $\forall s$
 $\Rightarrow h_1$ **dominates** h_2 and is **more efficient** for search.
- Theorem:** For any admissible heuristics h_1 and h_2 , define
 $h(s) = \max\{h_1(s), h_2(s)\}$
- $h(s)$ is admissible and dominates both h_1 and h_2 .
- '**Better**' heuristic = dominance = better search efficiency.
- Question:** Which one to choose from a collection of admissible heuristics h_1, \dots, h_m & none dominates any other?
- Answer:** $h(s) = \max\{h_1(s), \dots, h_m(s)\}$ dominates all the others.

取大的能防止过多搜索，提高效率

Quantify Search Efficiency

- Effective Branching Factor b^* :** For a solution from A*, calculate b^* satisfying: $N = b^* + (b^*)^2 + \dots + (b^*)^d$
 - N : #nodes of the solution,
 - d : depth of the solution tree.
 - E.g., A* finds a solution at depth 5 using 52 nodes $\Rightarrow b^* = 1.92$.
- Good heuristics have b^* close to 1 \Rightarrow large problems solved at reasonable computational cost.
- b^* quantifies search efficiency of heuristics.**

希望 b^* 越接近 1 越好

Empirical : Factor b^*

- Aim:** Compare h_1 and h_2 regarding the search efficiency.
- Setting:** Generate 1200 random problems with $d = \{2, \dots, 24\}$ and solve them with IDS and A* with h_1 & h_2 .
- Note:** IDS – a baseline.

d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

- h_2 is '**better**' than h_1 regarding search efficiency.
- This **goodness** is reflected by b^* being closer to 1.
- A^* with h_2 performs much better than IDS.

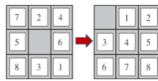
1.25

Generate Admissible Heuristics

① from Relaxed Problems

For 8-puzzle problem:

- **Real Rule:** A tile can only move to the **adjacent empty** square.
- **Relaxed rules:** h_{mis} and h_{1stp} are admissible
 - R1: A tile can move **anywhere** $\Rightarrow h_{mis}(s) = \#(\text{misplaced tiles})$.
 - R2: A tile can move one step in **any direction** regardless of an occupied neighbour $\Rightarrow h_{1stp}(s) = \#(\text{1-step move})$ to reach goal.
- **Optimal solutions to problems with R1, R2 are easier to find.**



Relaxed Problem

- **Relaxed problem:** a problem with **relaxed rules** on the action.
- E.g. 8-puzzle problems with R1 and R2.
- **Theorem:** The cost of an optimal solution to a **relaxed problem** is an **admissible heuristic** for the original problem.
- No wonder h_{mis} and h_{1stp} are admissible.

② from sub-problems

- **Subproblem**
 - **Task:** get tiles 1, 2, 3 and 4 into their correct positions.
 - **Relaxation:** move them disregarding the others.
- **Theory:** $\text{cost}^*(\text{subproblem}) < \text{cost}^*(\text{original})$.
 - $\text{cost}^*(\text{subproblem})$: the cost of the optimal solution of this subproblem.

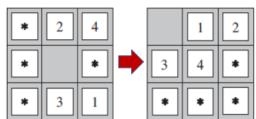


Fig.1. A subproblem of 8-puzzle.

Subproblem and Admissible Heuristics

- **Admissible $h_{sub}^*(s)$:** estimate the cost from s to the subproblem goal.
 - E.g. $h_{sub}^{(1,2,3,4)}$ is the cost to solve the 1-2-3-4 subproblem.
- **Theorem:** $h_{sub}(s)$ dominates $h_{1stp}(s)$.
 - $h_{sub}(s) = \max\{h_{sub}^{(1,2,3,4)}(s), h_{sub}^{(2,3,4,5)}(s), \dots\}$.

Disjoint Subproblems

- **Question:** Will the **addition of heuristics** from subproblem (1-2-3-4) and (5-6-7-8) give an **admissible heuristic**, considering the two subproblems are not overlapped?
- **Answer:** No, since they always **share some moves**.
- **Question:** What if **not count** those shared moves?
- **Answer:** $h_{sub}^{(1,2,3,4)}(s) + h_{sub}^{(5,6,7,8)}(s) \leq c^*(s) \Rightarrow \text{admissible}$.
 - Disjoint pattern database

③ from Experience

For 8-puzzle problem:

- Solve many 8-puzzles to obtain **many examples**.
- Each **example** consists of a state from the solution path and the actual cost of the solution from that point.
- These **examples** are our '**experience**' for this problem.
- **Question:** How to learn $h(s)$ from these **experience**?

Learn Heuristics from Experience

- **Question:** What are the **good experience features**?

- **Answer:** **Relevant** to predicting the states' cost to Goal, e.g.
 - $x_1(s)$: #(displaced tiles).
 - $x_2(s)$: #(pairs of adjacent tiles) that are not adjacent in Goal state.

- **Question:** How to learn h from those **relevant experience features**?

- **Answer:** (e.g.) Construct model as

$$h(s) = w_1x_1(s) + w_2x_2(s),$$

where w_1, w_2 are model parameters to learn from training data by a learning method such as neural networks and decision trees.

权值 W 即是经验生成的

Search Methods

- **Uniformed search:** Use **no** domain knowledge.

- BFS, DFS, DLS, IDS, UCS

- **Informed search:** Use a heuristic function that **estimates** how close a state is **to the goal**.

- Greedy search, A*, IDA*

We can organize the algorithms into pairs where the first proceeds by layers, and the other proceeds by subtrees.

(1) **Iterate on Node Depth:**

- BFS searches layers of increasing node depth.
- IDS searches subtrees of increasing node depth.

(2) **Iterate on Path Cost + Heuristic Function:**

- A* searches layers of increasing path cost + heuristic function.
- IDA* searches subtrees of increasing path cost + heuristic function.

Which cost function?

- UCS searches layers of increasing path cost.
- Greedy best first search searches layers of increasing heuristic function.
- A* search searches layers of increasing path cost + heuristic function.

Local Search

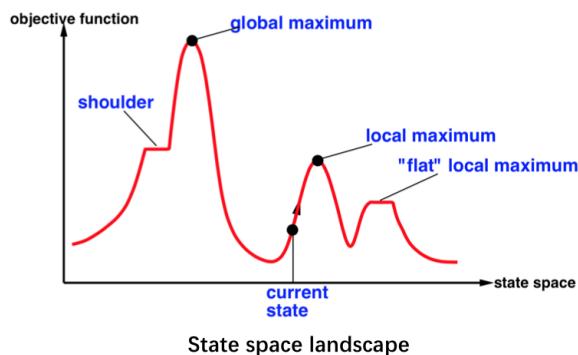
- Search algorithms seen so far are designed to **explore search spaces systematically**.
- Problems: observable, deterministic, known environments
- where the solution is a sequence of actions.
- Real-World problems are more complex.
- When a goal is found, the path to that goal constitutes a solution to the problem. But, depending on the applications, **the path may or may not matter**.
- If the path does not matter/systematic search is not possible, then consider another class of algorithms.
- In such cases, we can use iterative improvement algorithms, **Local search**.
- Also useful in pure **optimization problems** where the goal is to find the best state according to an **optimization function**.
- **Examples:**
 - Integrated circuit design, telecommunications network optimization, etc.
 - 8-queen: what matters is the final configuration of the puzzle, not the intermediary steps to reach it.

目标是找到最优状态，根据最优修改

- Idea: keep a single "current" state, and try to improve it.
- Move only to neighbors of that node.
- Advantages:
 - No need to maintain a search tree.
 - Use very little memory.
 - Can often find good enough solutions in continuous or large state spaces.

Local Search Algorithms:

- Hill climbing (steepest ascent/descent).
- Simulated Annealing: inspired by statistical physics.
- Local beam search.
- Genetic algorithms: inspired by evolutionary biology.



Hill climbing

- Also called greedy local search.
- Looks only to immediate good neighbors and not beyond.
- Search moves uphill: moves in the direction of increasing elevation/value to find the top of the mountain.
- Terminates when it reaches a **peak**.
- Can terminate with a local maximum, global maximum or can get stuck and no progress is possible.
- A node is a state and a value.

Pseudo-code

```
function HILL-CLIMBING(initialState)
  returns State that is a local maximum

  initialize current with initialState

  loop do
    neighbor = a highest-valued successor of current

    if neighbor.value ≤ current.value:
      return current.state

    current = neighbor
```

Variants

- Other variants of hill climbing include
- Sideways moves** escape from a plateau where best successor has same value as the current state.
 - Random-restart** hill climbing overcomes local maxima: keep trying! (either find a goal or get several possible solution and pick the max).
 - Stochastic** hill climbing chooses at random among the uphill moves.

只关注当前状态并优化它，只选择该状态附近的状态

爬山法
模拟退火
局部搜索
遗传算法

仅看最优邻居

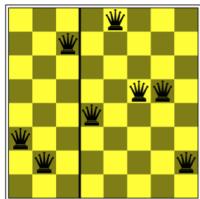
顺着使评估值增长的方向不断前进直到山顶。
当到达顶峰时停止。
容易停在局部最优

- Hill climbing** effective in general but depends on shape of the landscape. Successful in many real-problems after a reasonable number of restarts.
- Local beam search** maintains k states instead of one state. Select the k best successor, and useful information is passed among the states.
- Stochastic beam search** choose k successors are random. Helps alleviate the problem of the states agglomerating around the same part of the state space.

Stochastic: 做下一步的选择是随机的
beam: 批量选择

Genetic algorithms

- Genetic algorithm (GA) is a variant of stochastic beam search.
- Successor states are generated by combining two parents rather than modifying a single state.
- The process is inspired by natural selection.
- Starts with k randomly generated states, called population. Each state is an individual.
- An individual is usually represented by a string of 0's and 1's, or digits, a finite set.
- The objective function is called fitness function: better states have high values of fitness function.
- In the 8-queen problem, an individual can be represented by a string digits 1 to 8, that represents the position of the 8 queens in the 8 columns.



自然选择

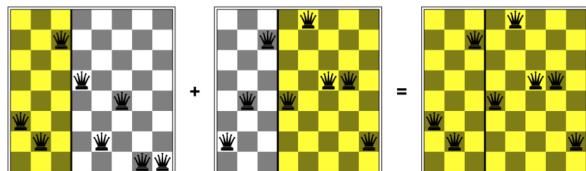
一开始有 k 个状态，即个体。个体合起来称物种群。
个体通常由0、1字符串来表示。

适应性函数：评估个体的优劣

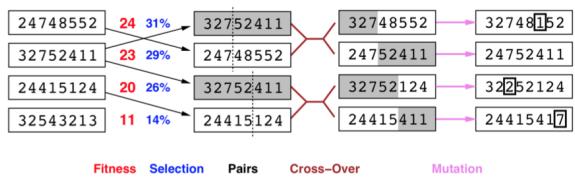
- The objective function is called fitness function: better states have high values of fitness function.
- Possible fitness function is the number of non-attacking pairs of queens.
- Fitness function of the solution: 28.
- Pairs of individuals are selected at random for reproduction w.r.t. some probabilities.
- A crossover point is chosen randomly in the string.
- Offspring are created by crossing the parents at the crossover point.
- Each element in the string is also subject to some mutation with a small probability.

41

交配产生后代时，双方均从自己的字符串中截取一部分生成新个体。交叉点是随机选择的。
后代会有一定概率产生突变



Generate successors from pairs of states.



Pseudo-code

```
function GENETIC-ALGORITHM(population, FITNESS-FN) returns an individual
  inputs: population, a set of individuals
    FITNESS-FN, a function that measures the fitness of an individual

  repeat
    new_population ← empty set
    for i = 1 to SIZE(population) do
      x ← RANDOM-SELECTION(population, FITNESS-FN)
      y ← RANDOM-SELECTION(population, FITNESS-FN)
      child ← REPRODUCE(x, y)
      if (small random probability) then child ← MUTATE(child)
      add child to new_population
    population ← new_population
  until some individual is fit enough, or enough time has elapsed
  return the best individual in population, according to FITNESS-FN

function REPRODUCE(x, y) returns an individual
  inputs: x, y, parent individuals

  n ← LENGTH(x); c ← random number from 1 to n
  return APPEND(SUBSTRING(x, 1, c), SUBSTRING(y, c + 1, n))
```

Simulated Annealing

- Hill Climbing → Simulated Annealing

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current ← problem.INITIAL
  while true do
    neighbor ← a highest-valued successor state of current
    if VALUE(neighbor) ≤ VALUE(current) then return current
    current ← neighbor

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  current ← problem.INITIAL
  for t = 1 to  $\infty$  do
    T ← schedule(t)
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E$  ← VALUE(current) - VALUE(next)
    if  $\Delta E < 0$  then current ← next //for maximization
    else current ← next only with probability  $e^{-\Delta E/T}$ 
```

爬山法是贪婪局部搜索，容易陷入局部最优 / 高原。

模拟退火是允许下山的爬山法。

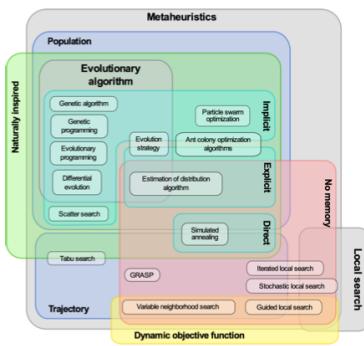
Further Studies on Metaheuristics

Comparison

Heuristic	Metaheuristic
• Problem-specific	• Problem-independent
• Can be used by a metaheuristic	• Can use different heuristics or a combination of heuristics

元启发算法是与问题无关的

Metaheuristics



Reproduction (crossover)
Mutation
“One general law, leading to the advancement of all organic beings, namely, multiply, vary, let the strongest live and weakest die.”
- Charles Darwin, *The Origin of Species*

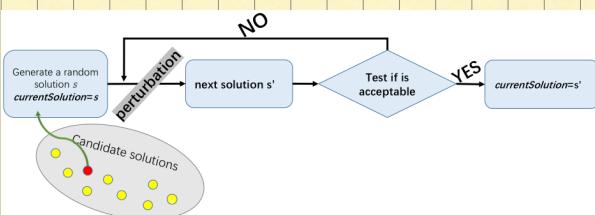
Evolutionary Computation

- It is the study of computational systems which use ideas and get inspirations from natural evolution.
- One of the principles borrowed is *survival of the fittest*.
- Evolutionary computation (EC) techniques can be used in optimisation, learning, and design.
- EC techniques do **not** require rich domain knowledge to use. However, domain knowledge can be incorporated into EC techniques.

适者生存
演化计算 (EC)

EC 算法无需大量的领域知识，但可以使用在 EC 算法中。

Generate-and-Test (G&T)

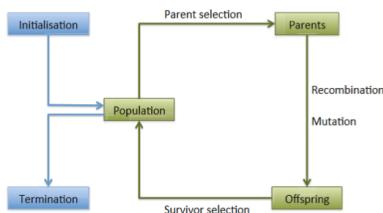


Steps

1. Generate the initial solution at random and denote it as the **current solution**.
2. Generate the **next solution** from the current one by **perturbation**.
3. Test whether the newly generated solution (**next solution**) is acceptable;
 1. Accepted it as the current solution **if yes**;
 2. Keep the current solution unchanged **otherwise**.
4. Go to Step 2 if the current solution is not satisfactory, stop otherwise.

EA: Population-based G & T

- **Generate:** Mutate and/or recombine individuals in a population.
- **Test:** Select the next generation from the parents and offspring.



A Simple Evolutionary Algorithm (EA)

- 1 Generate the initial population $P(0)$ at random
- 2 $i \leftarrow 0$ // Generation counter
- 3 WHILE halting criteria are not satisfied
- 4 Evaluate the fitness of each individual in $P(i)$
- 5 Select parents from $P(i)$ based on their fitness in $P(i)$
- 6 Generate offspring from the parents using crossover and mutation to form $P(i + 1)$
- 7 $i \leftarrow i + 1$

So how does this simple EA work?

Illustration Example

Let's use the simple EA with population size 4 to maximise the function

$$f(x) = x^2$$

with x in the integer interval $[0, 31]$, i.e., $x = 0, 1, \dots, 30, 31$.

- Population size = 4 \Leftrightarrow 4 individuals/chromosomes
- So, what is an **individual** or **chromosome**?

Example

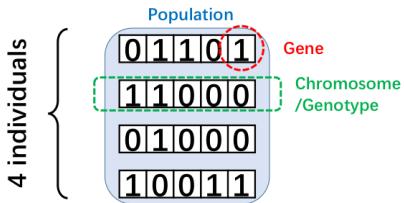
Encoding and decoding

- **Representation:** The first step of EA applications is **encoding** (i.e., the representation of chromosomes).
 - We adopt binary representation for integers.
 - 5 bits are used to represent integers up to 31.
 - Examples:

31	1	1	1	1	1	Encoding	Decoding	Evaluation
16	1	0	0	0	0	Encoding	Decoding	Evaluation

EA : Step 1

1. **Initialisation:** Generate initial population at random, e.g., 01101, 11000, 01000, 10011. These are **chromosomes** or **genotypes**.



EA : Step 2

2. **Evaluation:** Calculate fitness value for each individual.

- a) **Decode** the individual into an integer (called **phenotypes**):

$$01101 \rightarrow 13; 11000 \rightarrow 24, 01000 \rightarrow 8, 10011 \rightarrow 19;$$

- b) **Evaluate** the fitness according to $f(x) = x^2$:

$$f(13) = 169, f(24) = 576, f(8) = 64, f(19) = 361.$$

EA : Step 3-a

3. Crossover:

- a) Select two individuals for crossover based on their fitness. If **roulette-wheel selection** is used, then

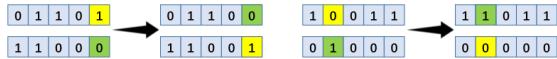
$$P_i = \frac{f_i}{\sum_j f_j}$$

Two offspring are often produced and added to an intermediate population. Repeat this step until the intermediate population is filled.
In our example:

$$P_1(13) = \frac{169}{1170} = 0.14, P_2(24) = \frac{576}{1170} = 0.49, P_3(8) = \frac{64}{1170} = 0.06, P_4(19) = \frac{361}{1170} = 0.31$$

EA : Step 3-b

- b) Examples of **crossover**



Now the intermediate population is 01100, 11011, 11011, 00000.

EA : Steps 4&5

4. Apply **mutation** to individuals in the intermediate population with a **small** probability. A simple mutation is bit-flipping. For example, we may have the following new population $P(1)$ after random mutation:

Example:

0	1	1	0	1	1	1	0	1	1
1	1	0	0	1	1	0	0	0	0

5. Go to step 2 if not stop.

Different Evolutionary Algorithms

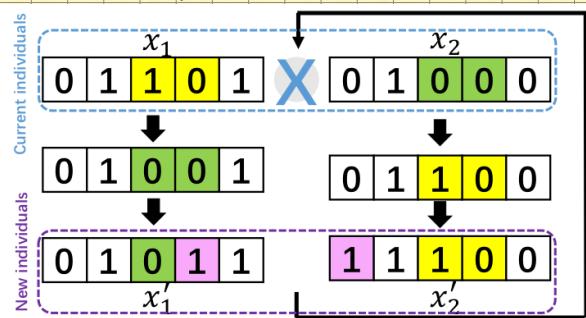
- There are several well-known EAs with different
 - historical backgrounds,
 - representations,
 - variation operators,
 - and selection schemes.

In fact, EAs refer to **a whole family of algorithms**, not a single algorithm.

EA families

- Genetic Algorithms (GAs)
- Evolutionary Programming (EP)
- Evolution Strategies (ES)
- Genetic Programming (GP)
- ...

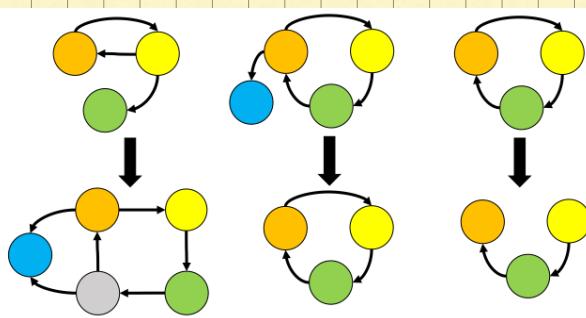
Genetic Algorithms (GAs)



- First formulated by Holland for adaptive search and by his students for optimisation from mid 1960s to mid 1970s.
- Binary strings have been used extensively as individuals (**chromosomes**).
- Simulate **Darwinian evolution**.
- Search operators are only applied to the **genotypic** representation (chromosome) of individuals.
- Emphasise the role of **recombination (crossover)**. Mutation is only used as a background operator.
- Often use **roulette-wheel** selection.

Representation	Bit-strings
Recombination	1-Point crossover
Mutation	Bit flip
Parent selection	Fitness proportional - implemented by Roulette Wheel
Survival selection	Generational

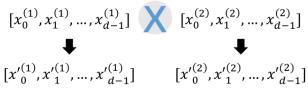
Evolutionary Programming (EP)



- First proposed by Fogel *et al.* in mid 1960s for simulating intelligence.
- Finite state machines** (FSMs) were used to represent individuals, although real-valued vectors have always been used in numerical optimisation.
- It is closer to **Larmackian evolution**.
- Search operators (mutations only) are applied to the **phenotypic** representation of individuals.
- It does **not** use any recombination.
- Usually use **tournament** selection.

有限状态机

Evolution Strategies (ES)

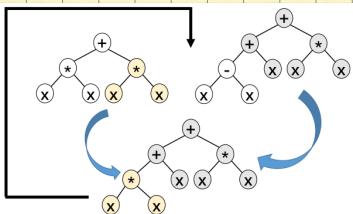


- First proposed by Rechenberg and Schwefel in mid 1960s for numerical optimisation.
- Real-valued vectors** are used to represent individuals.
- They are closer to **Larmackian evolution**.
- They do have recombination.
- They use **self-adaptive mutations**.

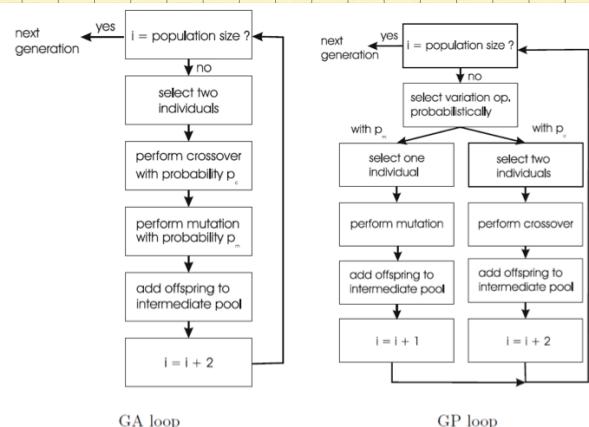
Representation	Real-valued vectors
Recombination	Discrete or intermediary
Mutation	Gaussian perturbation
Parent selection	Uniform random
Survivor selection	Deterministic elitist replacement by (μ, λ) or $(\mu + \lambda)$
Speciality	Self-adaptation of mutation step sizes

- μ : parents size
- λ : offspring size

Genetic Programming (GP)



- First used by de Garis to indicate the evolution of artificial neural networks, but used by Koza to indicate the evolution of computer programs.
- Trees** (especially Lisp expression trees) are often used to represent individuals.
- Both **crossover** and **mutation** are used.



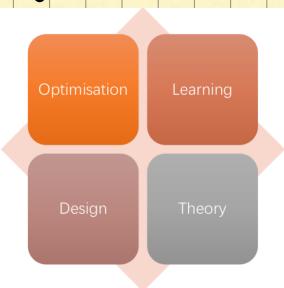
Preferred Term: Evolutionary Algorithms

- EAs face the same fundamental issues as those classical AI faces, i.e., representation, and search.
- Although GAs, EP, ES, and GP are different, they are all different variants of **population-based generate-and-test** algorithms. They share more similarities than differences!
- A better and more general term to use is evolutionary algorithms (EAs).

Variants in Operators

- **Crossover/Recombination:** one-point crossover, two-point crossover, uniform crossover, intermediate crossover, etc.
- **Mutation:** bit-flipping, Gaussian mutation, Cauchy mutation, etc.
- **Selection:** roulette wheel selection (fitness proportional selection), tournament selection, rank-based selection (linear and nonlinear), etc.
- **Replacement Strategy:** generational, steady-state (continuous), etc.
- **Specialised Operators:** multi-parent recombination, inversion, order-based crossover, etc.

Major Areas in EC



Evolutionary Optimization

- Numerical (global) optimisation.
- Combinatorial optimisation (of NP-hard problems).
- Mixed optimisation.
- Constrained optimisation.
- Multi-objective optimisation.
- Optimisation in a dynamic environment (with a dynamic fitness function).

Evolutionary Learning

Evolutionary learning can be used in supervised, unsupervised and reinforcement learning.

- Learning classifier systems (Rule-based systems).
- Evolutionary artificial neural networks.
- Evolutionary fuzzy logic systems.
- Co-evolutionary learning.

EC techniques are particularly good at exploring unconventional designs which are very difficult to obtain by hand.

- Evolutionary design of artificial neural networks.
- Evolutionary design of electronic circuits.
- Evolvable hardware.
- Evolutionary design of (building) architectures.

Evolutionary Design

Summary

- Evolutionary algorithms can be regarded as population-based generate-and-test algorithms.
- Evolutionary computation techniques can be used in optimisation, learning and design.
- Evolutionary computation techniques are flexible and robust.
- Evolutionary computation techniques are definitely useful tools in your toolbox, but there are problems for which other techniques might be more suitable.